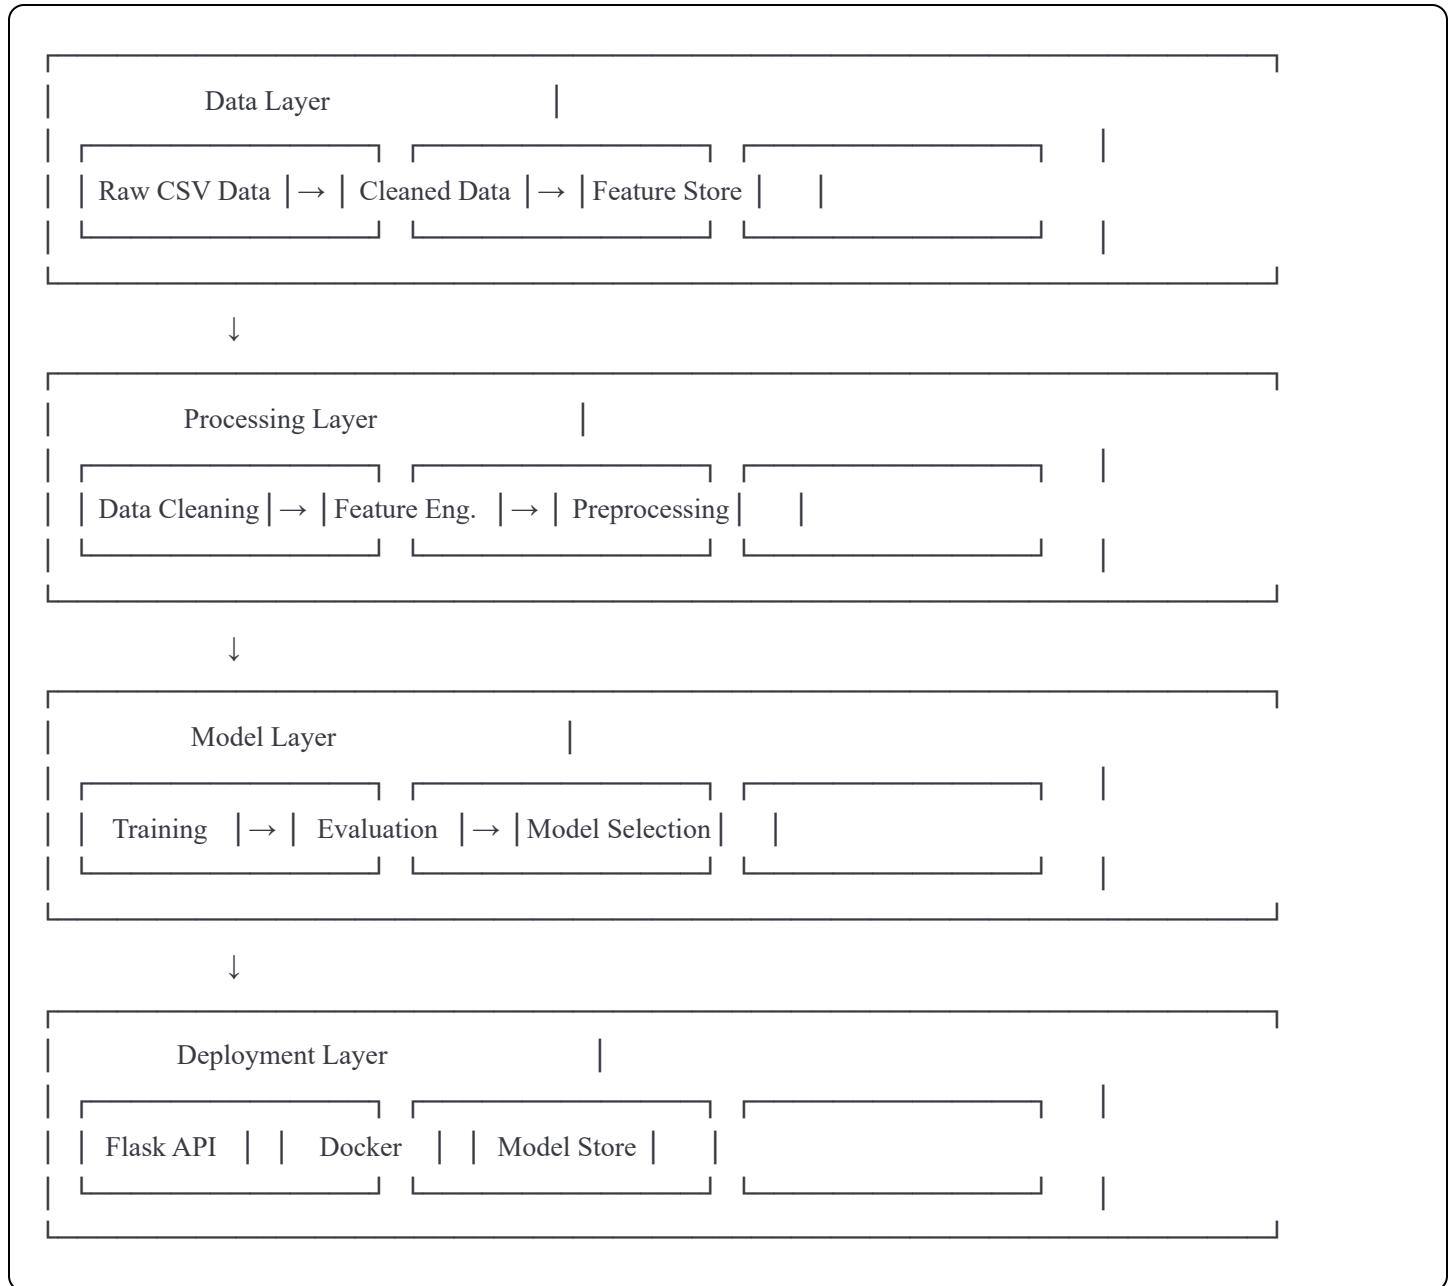


Heart Disease Prediction - Design Document

1. System Architecture

1.1 High-Level Architecture



1.2 Component Description

Data Layer

- **Raw Data Storage:** Original CSV file preservation
- **Cleaned Data:** Processed, validated data
- **Feature Store:** Engineered features ready for training

Processing Layer

- **Data Cleaning Module:** Missing value imputation, outlier detection
- **Feature Engineering:** Creating derived features, encoding
- **Preprocessing Pipeline:** Scaling, transformation, balancing

Model Layer

- **Training Engine:** Multiple algorithm training with cross-validation
- **Evaluation Module:** Performance metrics calculation
- **Model Selection:** Best model identification and optimization

Deployment Layer

- **REST API:** Flask-based prediction service
- **Containerization:** Docker for portability
- **Model Persistence:** Serialized model storage

2. Data Processing Design

2.1 Data Cleaning Strategy

Missing Value Handling

```
python
```

Strategy by Feature Type:

Numerical Features:

- └─ Age, BMI, Blood Pressure → Median imputation
- └─ Cholesterol, Triglycerides → Mean imputation (grouped by Age)
- └─ CRP, Homocysteine → KNN imputation (k=5)

Categorical Features:

- └─ Exercise Habits, Stress Level → Mode imputation
- └─ Alcohol Consumption → Special category "Unknown"
- └─ Binary features → Most frequent value

Outlier Detection

- **Method:** IQR (Interquartile Range) method
- **Action:** Cap at $1.5 * \text{IQR}$ boundaries

- **Features:** Blood Pressure, Cholesterol, BMI, Triglycerides

Data Validation Rules

```
python
```

Validation Rules:

```
|--- Age: 18 <= age <= 120
|--- Blood Pressure: 80 <= bp <= 250
|--- BMI: 10 <= bmi <= 60
|--- Cholesterol: 100 <= chol <= 500
└--- Sleep Hours: 0 <= sleep <= 24
```

2.2 Feature Engineering

Derived Features

1. Age Groups:

- Young Adult (18-35)
- Middle Age (36-55)
- Senior (56-80)

2. BMI Categories:

- Underweight (<18.5)
- Normal (18.5-24.9)
- Overweight (25-29.9)
- Obese (≥ 30)

3. Risk Score: Composite risk based on multiple factors

```
python
```

```
risk_score = (
    high_bp * 0.3 +
    diabetes * 0.25 +
    smoking * 0.2 +
    family_history * 0.15 +
    low_exercise * 0.1
)
```

4. Cholesterol Ratio: HDL/LDL interaction

5. Lifestyle Score: Exercise + Sleep + Stress combination

Feature Encoding Strategy

python

Encoding Approach:

Binary Features (Yes/No, Male/Female):

└─ Label Encoding (0, 1)

Ordinal Features (Low/Medium/High):

└─ Ordinal Encoding (0, 1, 2)

Nominal Features (if any):

└─ One-Hot Encoding

Target Variable:

└─ Label Encoding (No=0, Yes=1)

2.3 Feature Scaling

- **Method:** StandardScaler (zero mean, unit variance)
- **Applied to:** All numerical features
- **Pipeline:** Fit on training set, transform on test set

3. Model Design

3.1 Baseline Models

Model Selection Rationale

1. Logistic Regression

- Pros: Interpretable, fast, probability output
- Use: Baseline, feature importance
- Hyperparameters: C (regularization), penalty

2. Random Forest

- Pros: Handles non-linearity, feature importance
- Use: Ensemble baseline, robust to outliers
- Hyperparameters: n_estimators, max_depth, min_samples_split

3. Gradient Boosting (XGBoost)

- Pros: High performance, handles imbalance
- Use: Primary candidate for best performance
- Hyperparameters: learning_rate, max_depth, n_estimators

4. Support Vector Machine (SVM)

- Pros: Effective in high dimensions
- Use: Comparison benchmark
- Hyperparameters: C, kernel, gamma

5. LightGBM

- Pros: Fast training, efficient memory usage
- Use: Alternative to XGBoost
- Hyperparameters: num_leaves, learning_rate

3.2 Handling Class Imbalance

Techniques to Implement

1. SMOTE (Synthetic Minority Over-sampling Technique)

- Generate synthetic samples for minority class
- Apply only on training data

2. Class Weights

- Adjust loss function to penalize minority errors more
- Parameter: `class_weight='balanced'`

3. Stratified Sampling

- Maintain class distribution in train/test split
- Use `stratify` parameter

4. Ensemble Methods

- Inherently handle imbalance better
- Combine with sampling techniques

3.3 Hyperparameter Tuning

Search Strategy

python

Approach: RandomizedSearchCV

- Reason: Faster than GridSearch, good results
- Cross-Validation: 5-fold Stratified K-Fold
- Scoring: F1-Score (balance precision/recall)
- Iterations: 50 random combinations

Parameter Grids:

Random Forest:

- n_estimators: [100, 200, 300, 500]
- max_depth: [10, 20, 30, None]
- min_samples_split: [2, 5, 10]
- min_samples_leaf: [1, 2, 4]
- max_features: ['sqrt', 'log2']

XGBoost:

- learning_rate: [0.01, 0.05, 0.1, 0.2]
- max_depth: [3, 5, 7, 9]
- n_estimators: [100, 200, 300]
- subsample: [0.6, 0.8, 1.0]
- colsample_bytree: [0.6, 0.8, 1.0]
- scale_pos_weight: [1, 2, 4] (for imbalance)

3.4 Model Evaluation Design

Metrics Framework

python

Primary Metrics:

- └── F1-Score: Harmonic mean of precision/recall
- └── ROC-AUC: Area under ROC curve
- └── Precision-Recall AUC: For imbalanced data

Secondary Metrics:

- └── Accuracy: Overall correctness
- └── Precision: Positive predictive value
- └── Recall: Sensitivity
- └── Specificity: True negative rate
- └── Matthews Correlation Coefficient

Confusion Matrix Analysis:

- └── True Positives (TP): Correctly predicted disease
- └── True Negatives (TN): Correctly predicted healthy
- └── False Positives (FP): Incorrectly predicted disease
- └── False Negatives (FN): Missed disease cases

Cross-Validation Strategy

- **Method:** Stratified K-Fold ($k=5$)
- **Rationale:** Maintains class distribution in each fold
- **Reporting:** Mean \pm Standard Deviation for each metric

4. Code Architecture

4.1 Project Structure

```
heart_disease_prediction/
|
├── data/
│   ├── raw/          # Original data
│   ├── processed/    # Cleaned data
│   └── features/     # Engineered features
|
└── notebooks/
    ├── 01_EDA.ipynb   # Exploratory analysis
    ├── 02_Data_Cleaning.ipynb # Preprocessing
    ├── 03_Feature_Engineering.ipynb
    ├── 04_Model_Training.ipynb
    └── 05_Model_Evaluation.ipynb
```

```
src/
    __init__.py
    data/
        __init__.py
        loader.py      # Data loading
        cleaner.py    # Cleaning functions
        preprocessor.py # Preprocessing pipeline
    features/
        __init__.py
        engineering.py # Feature creation
        selection.py   # Feature selection
    models/
        __init__.py
        train.py       # Training logic
        evaluate.py    # Evaluation metrics
        predict.py     # Prediction interface
    visualization/
        __init__.py
        plots.py       # Plotting functions
    utils/
        __init__.py
        helpers.py     # Utility functions
models/
    trained/          # Saved models
    scalers/         # Preprocessing objects
    encoders/        # Encoding mappings
deployment/
    app.py           # Flask application
    Dockerfile       # Container definition
    requirements.txt # Dependencies
    config.py        # Configuration
tests/
    test_data.py     # Data tests
    test_models.py   # Model tests
    test_api.py      # API tests
docs/
```

```
|   └── PROJECT_SPECIFICATION.md  
|   └── DESIGN.md  
|   └── API_DOCUMENTATION.md  
|       └── MODEL_REPORT.md  
  
└── requirements.txt      # Python dependencies  
└── setup.py              # Package installation  
└── .gitignore            # Git ignore rules  
└── README.md             # Project overview
```

4.2 Module Design

Data Processing Module

```
python  
  
class DataCleaner:  
    def __init__(self, config):  
        self.config = config  
  
    def handle_missing_values(self, df):  
        """Impute missing values based on strategy"""  
        pass  
  
    def detect_outliers(self, df):  
        """Identify and handle outliers"""  
        pass  
  
    def validate_data(self, df):  
        """Ensure data meets validation rules"""  
        pass
```

Feature Engineering Module

```
python
```

```

class FeatureEngineer:
    def __init__(self):
        self.encoders = {}

    def create_derived_features(self, df):
        """Generate new features"""
        pass

    def encode_categorical(self, df):
        """Encode categorical variables"""
        pass

    def scale_features(self, df):
        """Standardize numerical features"""
        pass

```

Model Training Module

```

python

class ModelTrainer:
    def __init__(self, model_config):
        self.config = model_config
        self.models = {}

    def train_model(self, X, y, model_type):
        """Train specific model type"""
        pass

    def tune_hyperparameters(self, X, y, model):
        """Optimize hyperparameters"""
        pass

    def cross_validate(self, X, y, model):
        """Perform cross-validation"""
        pass

```

5. Deployment Design

5.1 API Design

Endpoints

POST /predict

 └─ Input: JSON with patient features

 └─ Output: Prediction + probability

 └─ Example:

 Request: {"age": 45, "gender": "Male", ...}

 Response: {"prediction": "Yes", "probability": 0.73}

GET /health

 └─ Purpose: Health check

 └─ Output: {"status": "healthy", "model_loaded": true}

GET /feature_importance

 └─ Purpose: Model interpretation

 └─ Output: Feature importance scores

POST /batch_predict

 └─ Purpose: Bulk predictions

 └─ Input: List of patient records

API Security

- Input validation
- Rate limiting
- HTTPS only (in production)
- API key authentication (optional)

5.2 Docker Configuration

Container Specs

dockerfile

Base Image: python:3.9-slim

Dependencies: requirements.txt

Exposed Port: 5000

Health Check: /health endpoint

Volume Mount: /models for persistence

5.3 Model Versioning

- Filename format: `model_v{version}_{date}.pkl`
- Metadata: Training date, metrics, features used

- Rollback capability: Keep previous versions

6. Testing Strategy

6.1 Unit Tests

- Data validation functions
- Feature engineering correctness
- Model loading/saving
- API endpoint responses

6.2 Integration Tests

- End-to-end pipeline
- API with model predictions
- Docker container functionality

6.3 Model Tests

- Prediction format validation
- Performance on validation set
- Handling edge cases

7. Monitoring and Logging

7.1 Logging Strategy

python

Logging Levels:

- └── DEBUG: Detailed diagnostic info
- └── INFO: General information flow
- └── WARNING: Something unexpected
- └── ERROR: Serious problems

Log Content:

- └── Data loading events
- └── Model training progress
- └── Prediction requests
- └── Error stack traces
- └── Performance metrics

7.2 Model Monitoring

- Track prediction distribution
- Monitor feature drift
- Log response times
- Alert on anomalies

8. Performance Optimization

8.1 Code Optimization

- Vectorized operations (NumPy)
- Efficient data structures
- Lazy loading for large datasets
- Parallel processing where applicable

8.2 Model Optimization

- Feature selection to reduce dimensionality
- Model compression techniques
- Batch prediction for efficiency
- Caching frequent predictions

Document Version: 1.0

Last Updated: January 2026

Author: ML Engineering Team